

ML Lab Report

Course: UE23CS352A: MACHINE LEARNING

Title: Artificial Neural Networks

Week-3

Name: Chandan Chatragadda	SRN: PES2UG23CS141	Section: 5C
---------------------------	--------------------	-------------

INTRODUCTION

1. Purpose of the lab:

The purpose of this lab was to gain hands-on experience in building a neural network from scratch for function approximation, without using high-level libraries like TensorFlow or PyTorch.

2. Tasks performed:

- Generated a synthetic polynomial dataset using my student ID using the already provided code.
- Implemented neural network components including ReLU activation, MSE loss, forward pass, and backpropagation.
- Split the data into 80% training and 20% testing and trained the neural network with Xavier initialization and gradient descent.
- Evaluated model performance using training/test loss, residuals, and prediction accuracy.

Dataset Description

1. Type of polynomial assigned:

Cubic polynomial: $y = 2.44x^3 - 0.47x^2 + 3.13x + 9.69$

2. Number of samples, features, noise level:

- **Samples:** 100,000 total (80,000 training, 20,000 testing).
- **Features:** 1 input feature (x), 1 output (y).
- **Noise:** Gaussian noise $e \sim N(0, 2.21)$.

Methodology


- A synthetic dataset was generated from the assigned polynomial with added Gaussian noise.
- Data was standardized using StandardScaler for both input and output variables.
- The neural network architecture used was: Input(1) → Hidden(32) → Hidden(72) → Output(1).
- Xavier initialization was applied to all weights, with biases set to zero.
- The forward pass applied linear transformation and ReLU activations, followed by a

linear output layer.


- Backpropagation was implemented manually to compute gradients and update weights via gradient descent.
- Training loop ran for 500 epochs with early stopping (patience = 10).
- Performance was evaluated using MSE loss, residual plots, predicted vs. actual plots, and R^2 score.

Results and Analysis




Exp1 (BaseLine)



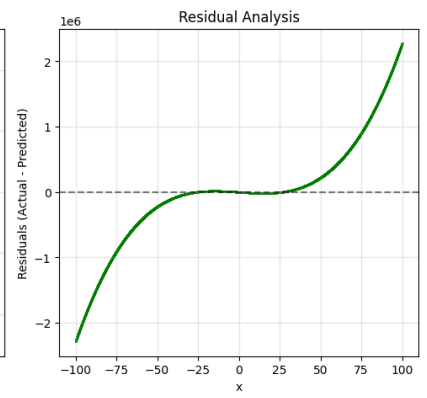


Exp2



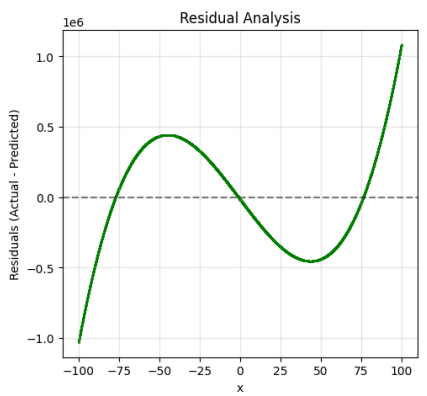

Exp 3



Exp 4



Exp 5



Experiment	Learning Rate	No. of epochs	Optimizers	Activation Function	Final Training Loss	Final Test Loss	R ² Score
Exp1 (Base)	0.005	500	Gradient Descent	ReLU	0.142655	0.138746	0.8575
Exp2	0.095	500	Gradient Descent	ReLU	0.001556	0.001481	0.9985
Exp3	0.095	350	Gradient Descent	ReLU	0.003228	0.003067	0.9968
Exp4	0.005	500	Gradient Descent	Sigmoid	0.824169	0.800634	0.1774
Exp5	0.075	750	Gradient Descent	Sigmoid	0.178014	0.173953	0.8213

Final Test MSE

From the table:

- The best performing model (Exp2) achieved a Final Test MSE of 0.001481 with an R² of 0.9985.
- This indicates the trained neural network almost perfectly approximated the target cubic polynomial with noise.

Base Model (Exp1 – ReLU, lr=0.005): The model learned reasonably well but had relatively higher loss and a modest R² of 0.8575. This suggests underfitting, as the network could not fully capture the cubic polynomial with a small learning rate.

Exp2 & Exp3 (ReLU, lr=0.095): Both runs achieved extremely low losses and very high R² (>0.996). The close match between training and test loss shows that the model generalized well and avoided overfitting. This configuration is the most effective for the dataset.

Exp4 (Sigmoid, lr=0.005): Performance degraded severely (R²=0.1774). The Sigmoid activation caused vanishing gradients, leading to underfitting where it failed to approximate the polynomial.

Exp5 (Sigmoid, lr=0.075): Performance improved compared to Exp4, but still lagged behind ReLU models. With $R^2=0.8213$, it still showed signs of underfitting, though it performed better due to the higher learning rate.

Conclusion

This lab demonstrated how to build a neural network from scratch to approximate a cubic polynomial. The results showed that ReLU activation with a higher learning rate gave the best performance, while Sigmoid often led to underfitting. The best model achieved a very low test MSE (0.001481) and an R^2 of 0.9985, showing it could predict the polynomial almost perfectly.